

Predicting Liver Disease:

The given dataset is related to Indian patients who have been tested for a liver disease. Based on chemical compounds (bilirubin, albumin, proteins, alkaline phosphatase) present in human body and tests like SGOT, SGPT the outcome mentioned is whether person is a patient i.e, whether he needs to be diagnosed further or not.

Objective:

Perform data cleansing, and required transformations and build a predictive model which will be able to predict most of the cases accurately.

Attributes:

Following are the feature names for the given data: Age, Gender, Total_Bilirubin, Direct_Bilirubin, Alkaline_Phosphotase, Alamine_Aminotransferase, Aspartate_Aminotransferase, Total_Protiens, Albumin, Albumin_and_Globulin_Ratio, Class.

Step 1:

Importing the libraries required -

```
In [54]: import pandas as pd
import numpy as np
import seaborn as sns
from sklearn.utils import resample
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
```

Step 2:

Loading the Data -

```
In [3]: data = pd.read_csv("C:/Users/Rohit/Downloads/IndianLiverPatientData.txt", delimiter="\t", header=None, index_col=0)
data.head()
```

Out[3]:

	1	2	3	4	5	6	7	8	9	10	11
0											
1	65	Female	0.7	0.1	187	16	18	6.8	3.3	0.90	No
2	62	Male	10.9	5.5	699	64	100	7.5	3.2	0.74	No
3	62	Male	7.3	4.1	490	60	68	7.0	3.3	0.89	No
4	58	Male	1.0	0.4	182	14	20	6.8	3.4	1.00	No
5	72	Male	3.9	2.0	195	27	59	7.3	2.4	0.40	No

As you can see we have no headings for the columns we give them manually

```
In [5]: data.columns = ["Age", "Gender", "Total_Bilirubin", "Direct_Bilirubin", "Alkaline_Phosphotase",
                        "Alamine_Aminotransferase", "Aspartate_Aminotransferase", "Total_Protiens", "Albumi
n",
                        , "Albumin_and_Globulin_Ratio", "Class"]
data.head()
```

Out[5]:

	Age	Gender	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Aminotransferase	Aspartate_Aminotransferase	Total_Pr
0								
1	65	Female	0.7	0.1	187	16	18	
2	62	Male	10.9	5.5	699	64	100	
3	62	Male	7.3	4.1	490	60	68	
4	58	Male	1.0	0.4	182	14	20	
5	72	Male	3.9	2.0	195	27	59	

We can check the summary by -

```
In [8]: summary = data.describe(include='all') # include all so that it takes categorical columns if there exist any  
print(summary)
```

	Age	Gender	Total_Bilirubin	Direct_Bilirubin \
count	583.000000	563	583.000000	583.000000
unique	NaN	2	NaN	NaN
top	NaN	Male	NaN	NaN
freq	NaN	421	NaN	NaN
mean	44.746141	NaN	3.298799	1.486106
std	16.189833	NaN	6.209522	2.808498
min	4.000000	NaN	0.400000	0.100000
25%	33.000000	NaN	0.800000	0.200000
50%	45.000000	NaN	1.000000	0.300000
75%	58.000000	NaN	2.600000	1.300000
max	90.000000	NaN	75.000000	19.700000

	Alkaline_Phosphotase	Alamine_Aminotransferase \
count	583.000000	583.000000
unique	NaN	NaN
top	NaN	NaN
freq	NaN	NaN
mean	290.576329	80.713551
std	242.937989	182.620356
min	63.000000	10.000000
25%	175.500000	23.000000
50%	208.000000	35.000000
75%	298.000000	60.500000
max	2110.000000	2000.000000

	Aspartate_Aminotransferase	Total_Protiens	Albumin \
count	583.000000	568.000000	583.000000
unique	NaN	NaN	NaN
top	NaN	NaN	NaN
freq	NaN	NaN	NaN
mean	109.910806	6.483979	3.141852
std	288.918529	1.084039	0.795519
min	10.000000	2.700000	0.900000
25%	25.000000	5.800000	2.600000
50%	42.000000	6.600000	3.100000
75%	87.000000	7.200000	3.800000
max	4929.000000	9.600000	5.500000

	Albumin_and_Globulin_Ratio	Class
count	579.000000	583
unique	NaN	2
top	NaN	No

freq	NaN	416
mean	0.947064	NaN
std	0.319592	NaN
min	0.300000	NaN
25%	0.700000	NaN
50%	0.930000	NaN
75%	1.100000	NaN
max	2.800000	NaN

Step 3:

Checking for null values -

```
In [9]: data.isnull().sum()
```

```
Out[9]: Age                0
Gender              20
Total_Bilirubin     0
Direct_Bilirubin    0
Alkaline_Phosphotase 0
Alamine_Aminotransferase 0
Aspartate_Aminotransferase 0
Total_Protiens      15
Albumin             0
Albumin_and_Globulin_Ratio 4
Class               0
dtype: int64
```

```
In [10]: # creating a copy of data if we mess something up
data1 = data.copy()
```

```
In [12]: data1['Gender'].value_counts()
```

```
Out[12]: Male      421
Female    142
Name: Gender, dtype: int64
```

```
In [13]: # as we have more values for "Male" we fill the missing values in Gender column with "Male"
data1['Gender'].fillna(data1['Gender'].mode()[0], inplace=True)
```

```
In [14]: # and we do the following for Albumin_and_Globulin_Ratio, Total_Protiens
for value in ['Albumin_and_Globulin_Ratio', 'Total_Protiens']:
    data1[value].fillna(data1[value].median(), inplace=True)
```

```
In [15]: # to check if NAs still exist
data1.isnull().sum()
```

```
Out[15]: Age                0
Gender                0
Total_Bilirubin       0
Direct_Bilirubin      0
Alkaline_Phosphotase  0
Alamine_Aminotransferase  0
Aspartate_Aminotransferase  0
Total_Protiens        0
Albumin               0
Albumin_and_Globulin_Ratio  0
Class                 0
dtype: int64
```

Step 4:

Performing Label Encoding on categorical columns -

```
In [16]: col_name = ['Gender', 'Class']
from sklearn import preprocessing
le={}
for value in col_name:
    le[value]=preprocessing.LabelEncoder()
for value in col_name:
    data1[value] = le[value].fit_transform(data1[value])

data1.info() # see if all values are either float or int
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 583 entries, 1 to 583
Data columns (total 11 columns):
Age                    583 non-null int64
Gender                 583 non-null int32
Total_Bilirubin        583 non-null float64
Direct_Bilirubin       583 non-null float64
Alkaline_Phosphotase   583 non-null int64
Alamine_Aminotransferase 583 non-null int64
Aspartate_Aminotransferase 583 non-null int64
Total_Protiens         583 non-null float64
Albumin               583 non-null float64
Albumin_and_Globulin_Ratio 583 non-null float64
Class                 583 non-null int32
dtypes: float64(5), int32(2), int64(4)
memory usage: 50.1 KB
```

Step 5:

To check for Outliers -

```
In [19]: summary2 = data1.describe(np.arange(0.9, 0.99, 0.01), include='all')  
print(summary2)
```


	Age	Gender	Total_Bilirubin	Direct_Bilirubin \
count	583.000000	583.000000	583.000000	583.000000
mean	44.746141	0.756432	3.298799	1.486106
std	16.189833	0.429603	6.209522	2.808498
min	4.000000	0.000000	0.400000	0.100000
50%	45.000000	1.000000	1.000000	0.300000
90%	66.000000	1.000000	7.860000	4.080000
91%	66.000000	1.000000	8.900000	4.500000
92%	67.440000	1.000000	10.900000	5.144000
93%	69.260000	1.000000	12.256000	6.252000
94%	70.000000	1.000000	14.816000	7.608000
95%	72.000000	1.000000	16.350000	8.400000
96%	72.720000	1.000000	18.288000	8.972000
97%	74.540000	1.000000	20.108000	10.108000
98%	75.000000	1.000000	23.072000	11.736000
max	90.000000	1.000000	75.000000	19.700000

	Alkaline_Phosphotase	Alamine_Aminotransferase \
count	583.000000	583.000000
mean	290.576329	80.713551
std	242.937989	182.620356
min	63.000000	10.000000
50%	208.000000	35.000000
90%	511.400000	140.000000
91%	549.440000	153.240000
92%	574.440000	166.880000
93%	610.520000	183.080000
94%	661.240000	198.560000
95%	698.100000	232.000000
96%	792.480000	341.440000
97%	908.560000	409.700000
98%	1103.600000	549.680000
max	2110.000000	2000.000000

	Aspartate_Aminotransferase	Total_Protiens	Albumin \
count	583.000000	583.000000	583.000000
mean	109.910806	6.486964	3.141852
std	288.918529	1.070136	0.795519
min	10.000000	2.700000	0.900000
50%	42.000000	6.600000	3.100000
90%	190.000000	7.900000	4.100000
91%	230.620000	7.900000	4.200000
92%	239.960000	8.000000	4.200000

93%	277.600000	8.000000	4.300000
94%	351.360000	8.000000	4.300000
95%	400.900000	8.100000	4.390000
96%	507.920000	8.200000	4.400000
97%	626.780000	8.300000	4.500000
98%	846.160000	8.500000	4.636000
max	4929.000000	9.600000	5.500000

	Albumin_and_Globulin_Ratio	Class
count	583.000000	583.000000
mean	0.946947	0.286449
std	0.318495	0.452490
min	0.300000	0.000000
50%	0.930000	0.000000
90%	1.300000	1.000000
91%	1.380000	1.000000
92%	1.400000	1.000000
93%	1.400000	1.000000
94%	1.400000	1.000000
95%	1.500000	1.000000
96%	1.507200	1.000000
97%	1.600000	1.000000
98%	1.700000	1.000000
max	2.800000	1.000000

```
In [20]: data2 = data1.loc[data1["Alkaline_Phosphotase"] < 909]
# here we took those samples where the specified column had value below 97th %ile

data3 = data2.loc[data2["Alamine_Aminotransferase"] < 500]
# here we took those samples where the specified column had value below 98th %ile

# we took this values because they seemed to have no outliers beyond those values
```

Step 6:

Checking multicollinearity and removing it -

```
In [23]: # independent variables
x = data3[['Age', 'Gender', 'Total_Bilirubin', 'Direct_Bilirubin',
          'Alkaline_Phosphotase', 'Alamine_Aminotransferase',
          'Aspartate_Aminotransferase', 'Total_Protiens', 'Albumin',
          'Albumin_and_Globulin_Ratio']]

# dependent variables
y = data3[['Class']]

corr2 = x.corr(method='pearson')
corr2
```

Out[23]:

	Age	Gender	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Aminotransferase	As
Age	1.000000	0.063305	0.020704	0.012033	0.009075	-0.053116	
Gender	0.063305	1.000000	0.107551	0.121324	0.054975	0.117867	
Total_Bilirubin	0.020704	0.107551	1.000000	0.860373	0.184910	0.223902	
Direct_Bilirubin	0.012033	0.121324	0.860373	1.000000	0.209227	0.264239	
Alkaline_Phosphotase	0.009075	0.054975	0.184910	0.209227	1.000000	0.245415	
Alamine_Aminotransferase	-0.053116	0.117867	0.223902	0.264239	0.245415	1.000000	
Aspartate_Aminotransferase	-0.007257	0.124726	0.294169	0.328998	0.156771	0.774318	
Total_Protiens	-0.185952	-0.096330	0.006380	0.012968	0.024365	0.067497	
Albumin	-0.265396	-0.114106	-0.224084	-0.234582	-0.136099	-0.001841	
Albumin_and_Globulin_Ratio	-0.207842	-0.024835	-0.200034	-0.194982	-0.224555	-0.075860	

In [25]: *# We remove albumin as it is correlated to both 'total protein' and 'albumin and globulin ratio'*

and we check the multicollinearity again

```
x = data3[['Age', 'Gender', 'Total_Bilirubin', 'Direct_Bilirubin',
           'Alkaline_Phosphotase', 'Alamine_Aminotransferase',
           'Aspartate_Aminotransferase', 'Total_Protiens',
           'Albumin_and_Globulin_Ratio']]
```

```
cor3 = x.corr(method = "pearson")
```

```
cor3
```

Out[25]:

	Age	Gender	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Aminotransferase	As
Age	1.000000	0.063305	0.020704	0.012033	0.009075	-0.053116	
Gender	0.063305	1.000000	0.107551	0.121324	0.054975	0.117867	
Total_Bilirubin	0.020704	0.107551	1.000000	0.860373	0.184910	0.223902	
Direct_Bilirubin	0.012033	0.121324	0.860373	1.000000	0.209227	0.264239	
Alkaline_Phosphotase	0.009075	0.054975	0.184910	0.209227	1.000000	0.245415	
Alamine_Aminotransferase	-0.053116	0.117867	0.223902	0.264239	0.245415	1.000000	
Aspartate_Aminotransferase	-0.007257	0.124726	0.294169	0.328998	0.156771	0.774318	
Total_Protiens	-0.185952	-0.096330	0.006380	0.012968	0.024365	0.067497	
Albumin_and_Globulin_Ratio	-0.207842	-0.024835	-0.200034	-0.194982	-0.224555	-0.075860	



In [26]: *# 'Alamine_Aminotransferase' is more specific to liver than 'Aspartate_Aminotransferase' so we keep
'Alamine_Aminotransferase' and drop 'Aspartate_Aminotransferase'*

```
x = data3[['Age', 'Gender', 'Total_Bilirubin', 'Direct_Bilirubin',
           'Alkaline_Phosphotase', 'Alamine_Aminotransferase',
           'Total_Protiens', 'Albumin_and_Globulin_Ratio']]
cor4 = x.corr(method="pearson")
cor4
```

Out[26]:


	Age	Gender	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Aminotransferase	Tot
Age	1.000000	0.063305	0.020704	0.012033	0.009075	-0.053116	
Gender	0.063305	1.000000	0.107551	0.121324	0.054975	0.117867	
Total_Bilirubin	0.020704	0.107551	1.000000	0.860373	0.184910	0.223902	
Direct_Bilirubin	0.012033	0.121324	0.860373	1.000000	0.209227	0.264239	
Alkaline_Phosphotase	0.009075	0.054975	0.184910	0.209227	1.000000	0.245415	
Alamine_Aminotransferase	-0.053116	0.117867	0.223902	0.264239	0.245415	1.000000	
Total_Protiens	-0.185952	-0.096330	0.006380	0.012968	0.024365	0.067497	
Albumin_and_Globulin_Ratio	-0.207842	-0.024835	-0.200034	-0.194982	-0.224555	-0.075860	

In [27]: *# As 'Total bilirubin' considers both water and fat soluble substances we drop 'direct bilirubin'*

```
x = data3[['Age', 'Gender', 'Total_Bilirubin',
           'Alkaline_Phosphotase', 'Alamine_Aminotransferase',
           'Total_Protiens', 'Albumin_and_Globulin_Ratio']]
cor5 = x.corr(method="pearson")
```

In [28]: *# we check the same with the help of heat map whether removing the above columns helped in reducing multicollinearity*
 sns.heatmap(coor5, vmin=-1, vmax=1, annot=True)

Out[28]: <matplotlib.axes._subplots.AxesSubplot at 0x123c6357f08>



Step 7:

Resampling (*Class Balancing*)

In [29]: data3['Class'].value_counts()

Out[29]: 0 388
 1 166
 Name: Class, dtype: int64

We observe here that our dependent variable has class imbalanced as class 0 has 388 observations and class 1 has 166 observations. Hence to resolve this issue we upsample the class 1 to class 0 so that it matches perfectly and doesn't affect our predictions.

```
In [31]: data4 = pd.concat([x, y], axis=1)
df_major = data4[data4['Class'] == 0]
df_minor = data4[data4['Class'] == 1]

upSample = resample(df_minor, replace=True, n_samples=388, random_state=0)

data5 = pd.concat([df_major, upSample])
data5['Class'].value_counts()
```

```
Out[31]: 1    388
         0    388
         Name: Class, dtype: int64
```

As you can see we have matched both the classes

Step 8:

Scaling the data

```
In [32]: x = data5[['Age', 'Gender', 'Total_Bilirubin',
                  'Alkaline_Phosphotase', 'Alamine_Aminotransferase',
                  'Total_Protiens', 'Albumin_and_Globulin_Ratio']]
y = data5['Class']

from sklearn.preprocessing import StandardScaler
scale = StandardScaler()
scale.fit(x)
x = scale.transform(x)
```

Step 9:

Splitting the data into training and testing

```
In [33]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=0)
```

Step 10:

Performing Logistic Regression

```
In [35]: model_lr = LogisticRegression()
model_lr.fit(x_train, y_train)

y_pred = model_lr.predict(x_test)
```

```
In [36]: # confusion matrix and some othe measures
cfm = confusion_matrix(y_test, y_pred)

print(cfm)
print(classification_report(y_test, y_pred))
print(accuracy_score(y_test, y_pred))
```

```
[[41 30]
 [19 66]]

      precision    recall  f1-score   support

     0       0.68      0.58      0.63        71
     1       0.69      0.78      0.73        85

 accuracy          0.69          156
 macro avg       0.69      0.68      0.68          156
 weighted avg    0.69      0.69      0.68          156

0.6858974358974359
```


The catch here is that while dividing the classes as the probability of being 1, the GLM by default assumes the threshold as 0.5 Hence we need to decide which threshold suits better by checking for errors manually.

```
In [37]: # this gives us a data frame for probabiltiy of each observation being 0 or 1  
y_pred_prob = model_lr.predict_proba(x_test)  
print(y_pred_prob) # first column is prob. of being 0 and 2nd column is prob. of beinf 1
```

```
[ [3.79015689e-01 6.20984311e-01]
[3.15061664e-01 6.84938336e-01]
[5.03658284e-01 4.96341716e-01]
[8.76892943e-01 1.23107057e-01]
[3.85386419e-01 6.14613581e-01]
[3.42904387e-01 6.57095613e-01]
[3.10812823e-01 6.89187177e-01]
[3.84841870e-01 6.15158130e-01]
[2.73809565e-01 7.26190435e-01]
[3.88145628e-01 6.11854372e-01]
[3.59816573e-01 6.40183427e-01]
[3.36302097e-01 6.63697903e-01]
[7.16060973e-01 2.83939027e-01]
[3.10763336e-01 6.89236664e-01]
[3.07018224e-01 6.92981776e-01]
[3.88145628e-01 6.11854372e-01]
[3.94509746e-01 6.05490254e-01]
[5.69212388e-01 4.30787612e-01]
[5.98679316e-01 4.01320684e-01]
[5.93751541e-01 4.06248459e-01]
[4.03036778e-01 5.96963222e-01]
[2.20803501e-01 7.79196499e-01]
[2.5669791e-01 7.43330209e-01]
[2.73809565e-01 7.26190435e-01]
[9.92174127e-01 7.82587320e-03]
[3.90471903e-01 6.09528097e-01]
[6.05807637e-01 3.94192363e-01]
[2.43451824e-01 7.56548176e-01]
[4.95418788e-01 5.04581212e-01]
[5.69212388e-01 4.30787612e-01]
[7.07644299e-01 2.92355701e-01]
[9.06688507e-01 9.33114931e-02]
[3.19267349e-01 6.80732651e-01]
[4.13974047e-01 5.86025953e-01]
[3.10812823e-01 6.89187177e-01]
[3.81609614e-01 6.18390386e-01]
[4.64449959e-01 5.35550041e-01]
[5.98568958e-01 4.01431042e-01]
[9.91743765e-01 8.25623548e-03]
[4.86507424e-01 5.13492576e-01]
[2.38835129e-01 7.61164871e-01]
[4.87346008e-01 5.12653992e-01]
[2.38397837e-01 7.61602163e-01]
```

[4.03036778e-01 5.96963222e-01]
[5.62671484e-01 4.37328516e-01]
[9.98515433e-01 1.48456746e-03]
[3.03010949e-01 6.96989051e-01]
[3.79014877e-01 6.20985123e-01]
[5.42738746e-01 4.57261254e-01]
[8.73041399e-01 1.26958601e-01]
[6.34869659e-01 3.65130341e-01]
[6.85270531e-01 3.14729469e-01]
[4.53845267e-01 5.46154733e-01]
[4.00771570e-01 5.99228430e-01]
[7.86112990e-01 2.13887010e-01]
[9.99996035e-01 3.96501846e-06]
[5.67778258e-01 4.32221742e-01]
[9.96198740e-01 3.80125971e-03]
[4.83863928e-01 5.16136072e-01]
[5.00577378e-01 4.99422622e-01]
[4.14810045e-01 5.85189955e-01]
[2.32527928e-01 7.67472072e-01]
[4.54921707e-01 5.45078293e-01]
[7.60498846e-01 2.39501154e-01]
[4.06605820e-01 5.93394180e-01]
[5.21287336e-01 4.78712664e-01]
[9.96963944e-01 3.03605569e-03]
[2.50372469e-01 7.49627531e-01]
[3.88145628e-01 6.11854372e-01]
[9.18223301e-01 8.17766993e-02]
[3.37235080e-01 6.62764920e-01]
[2.75462494e-01 7.24537506e-01]
[3.41451995e-01 6.58548005e-01]
[2.93762536e-01 7.06237464e-01]
[2.94212802e-01 7.05787198e-01]
[3.13506319e-01 6.86493681e-01]
[5.61831256e-01 4.38168744e-01]
[2.56669791e-01 7.43330209e-01]
[6.71117047e-01 3.28882953e-01]
[3.56325539e-01 6.43674461e-01]
[3.24748153e-01 6.75251847e-01]
[2.43451824e-01 7.56548176e-01]
[4.31327420e-01 5.68672580e-01]
[3.35251797e-01 6.64748203e-01]
[3.72372375e-01 6.27627625e-01]
[9.98919167e-01 1.08083282e-03]

[9.15354106e-01 8.46458940e-02]
[2.95017267e-01 7.04982733e-01]
[3.69020642e-01 6.30979358e-01]
[9.81787406e-01 1.82125938e-02]
[3.17188750e-01 6.82811250e-01]
[3.33774386e-01 6.66225614e-01]
[3.79014877e-01 6.20985123e-01]
[3.61134351e-01 6.38865649e-01]
[6.80960233e-01 3.19039767e-01]
[3.40114822e-01 6.59885178e-01]
[4.26154753e-01 5.73845247e-01]
[5.42738746e-01 4.57261254e-01]
[4.35508239e-01 5.64491761e-01]
[3.28734792e-01 6.71265208e-01]
[4.86507424e-01 5.13492576e-01]
[3.26824059e-01 6.73175941e-01]
[2.15863680e-01 7.84136320e-01]
[3.33575970e-01 6.66424030e-01]
[8.02320224e-01 1.97679776e-01]
[5.27855795e-01 4.72144205e-01]
[9.69709055e-01 3.02909451e-02]
[5.18770701e-01 4.81229299e-01]
[2.49224154e-01 7.50775846e-01]
[1.52992519e-01 8.47007481e-01]
[5.90766852e-01 4.09233148e-01]
[2.88823981e-01 7.11176019e-01]
[9.70865428e-01 2.91345724e-02]
[6.31984730e-01 3.68015270e-01]
[3.24748153e-01 6.75251847e-01]
[9.99702595e-01 2.97405135e-04]
[2.75462494e-01 7.24537506e-01]
[3.67938666e-01 6.32061334e-01]
[2.73520065e-01 7.26479935e-01]
[2.67506832e-01 7.32493168e-01]
[9.03930755e-01 9.60692452e-02]
[3.73131893e-01 6.26868107e-01]
[2.24104415e-01 7.75895585e-01]
[4.35525277e-01 5.64474723e-01]
[9.91571506e-01 8.42849427e-03]
[9.99615116e-01 3.84884044e-04]
[5.55836139e-01 4.44163861e-01]
[7.13179984e-01 2.86820016e-01]
[2.94107841e-01 7.05892159e-01]

```
[4.51388281e-01 5.48611719e-01]  
[4.34263767e-01 5.65736233e-01]  
[4.53845267e-01 5.46154733e-01]  
[4.37018484e-01 5.62981516e-01]  
[2.16875029e-01 7.83124971e-01]  
[5.23328064e-01 4.76671936e-01]  
[4.16002870e-01 5.83997130e-01]  
[9.88541960e-01 1.14580400e-02]  
[2.95017267e-01 7.04982733e-01]  
[6.05229494e-01 3.94770506e-01]  
[2.50372469e-01 7.49627531e-01]  
[5.30404301e-01 4.69595699e-01]  
[5.16963147e-01 4.83036853e-01]  
[9.86161563e-01 1.38384374e-02]  
[3.13506319e-01 6.86493681e-01]  
[2.58742733e-01 7.41257267e-01]  
[3.16627167e-01 6.83372833e-01]  
[9.99977898e-01 2.21019831e-05]  
[9.81575573e-01 1.84244272e-02]  
[2.32527928e-01 7.67472072e-01]  
[8.69552135e-01 1.30447865e-01]  
[7.69999199e-01 2.30000801e-01]  
[2.43451824e-01 7.56548176e-01]  
[3.85386419e-01 6.14613581e-01]  
[9.16525795e-01 8.34742052e-02]  
[2.75227144e-01 7.24772856e-01]  
[8.08347056e-01 1.91652944e-01]]
```

```
In [38]: # Now we check threshold wise which gives us Less Type 2 error and significantly acceptable Total Error and Type 1 error
for a in np.arange(0.2, 0.8, 0.01):
    predict_mine = np.where(y_pred_prob[:, 1] > a, 1, 0)
    cfm = confusion_matrix(y_test, predict_mine)
    total_err = cfm[0, 1]+cfm[1, 0]
    print("Errors at threshold ", a, ":", total_err, " , type 2 error :",
          cfm[1, 0], " , type 1 error:", cfm[0, 1])
```

```

Errors at threshold 0.2 : 46 , type 2 error : 1 , type 1 error: 45
Errors at threshold 0.2100000000000002 : 46 , type 2 error : 1 , type 1 error: 45
Errors at threshold 0.2200000000000003 : 45 , type 2 error : 1 , type 1 error: 44
Errors at threshold 0.2300000000000004 : 45 , type 2 error : 1 , type 1 error: 44
Errors at threshold 0.2400000000000005 : 43 , type 2 error : 1 , type 1 error: 42
Errors at threshold 0.2500000000000006 : 43 , type 2 error : 1 , type 1 error: 42
Errors at threshold 0.2600000000000006 : 43 , type 2 error : 1 , type 1 error: 42
Errors at threshold 0.2700000000000001 : 43 , type 2 error : 1 , type 1 error: 42
Errors at threshold 0.2800000000000001 : 43 , type 2 error : 1 , type 1 error: 42
Errors at threshold 0.2900000000000001 : 43 , type 2 error : 2 , type 1 error: 41
Errors at threshold 0.3000000000000001 : 44 , type 2 error : 3 , type 1 error: 41
Errors at threshold 0.3100000000000001 : 44 , type 2 error : 3 , type 1 error: 41
Errors at threshold 0.3200000000000001 : 46 , type 2 error : 5 , type 1 error: 41
Errors at threshold 0.3300000000000001 : 47 , type 2 error : 6 , type 1 error: 41
Errors at threshold 0.34000000000000014 : 47 , type 2 error : 6 , type 1 error: 41
Errors at threshold 0.35000000000000014 : 47 , type 2 error : 6 , type 1 error: 41
Errors at threshold 0.36000000000000015 : 47 , type 2 error : 6 , type 1 error: 41
Errors at threshold 0.37000000000000016 : 47 , type 2 error : 7 , type 1 error: 40
Errors at threshold 0.38000000000000017 : 47 , type 2 error : 7 , type 1 error: 40
Errors at threshold 0.3900000000000002 : 47 , type 2 error : 7 , type 1 error: 40
Errors at threshold 0.4000000000000002 : 47 , type 2 error : 8 , type 1 error: 39
Errors at threshold 0.4100000000000002 : 49 , type 2 error : 11 , type 1 error: 38
Errors at threshold 0.4200000000000002 : 49 , type 2 error : 11 , type 1 error: 38
Errors at threshold 0.4300000000000002 : 49 , type 2 error : 11 , type 1 error: 38
Errors at threshold 0.4400000000000002 : 48 , type 2 error : 13 , type 1 error: 35
Errors at threshold 0.45000000000000023 : 47 , type 2 error : 13 , type 1 error: 34
Errors at threshold 0.46000000000000024 : 49 , type 2 error : 15 , type 1 error: 34
Errors at threshold 0.47000000000000025 : 48 , type 2 error : 15 , type 1 error: 33
Errors at threshold 0.48000000000000026 : 47 , type 2 error : 16 , type 1 error: 31
Errors at threshold 0.49000000000000027 : 47 , type 2 error : 17 , type 1 error: 30
Errors at threshold 0.5000000000000002 : 49 , type 2 error : 19 , type 1 error: 30
Errors at threshold 0.5100000000000002 : 48 , type 2 error : 19 , type 1 error: 29
Errors at threshold 0.5200000000000002 : 48 , type 2 error : 21 , type 1 error: 27
Errors at threshold 0.5300000000000002 : 48 , type 2 error : 21 , type 1 error: 27
Errors at threshold 0.5400000000000003 : 49 , type 2 error : 22 , type 1 error: 27
Errors at threshold 0.5500000000000003 : 51 , type 2 error : 25 , type 1 error: 26
Errors at threshold 0.5600000000000003 : 51 , type 2 error : 25 , type 1 error: 26
Errors at threshold 0.5700000000000003 : 52 , type 2 error : 28 , type 1 error: 24
Errors at threshold 0.5800000000000003 : 53 , type 2 error : 29 , type 1 error: 24
Errors at threshold 0.5900000000000003 : 52 , type 2 error : 30 , type 1 error: 22
Errors at threshold 0.6000000000000003 : 56 , type 2 error : 34 , type 1 error: 22
Errors at threshold 0.6100000000000003 : 58 , type 2 error : 36 , type 1 error: 22
Errors at threshold 0.6200000000000003 : 63 , type 2 error : 42 , type 1 error: 21

```


Errors at threshold	0.6300000000000003	: 62	, type 2 error	: 44	, type 1 error	: 18
Errors at threshold	0.6400000000000003	: 63	, type 2 error	: 46	, type 1 error	: 17
Errors at threshold	0.6500000000000004	: 63	, type 2 error	: 47	, type 1 error	: 16
Errors at threshold	0.6600000000000004	: 62	, type 2 error	: 48	, type 1 error	: 14
Errors at threshold	0.6700000000000004	: 61	, type 2 error	: 50	, type 1 error	: 11
Errors at threshold	0.6800000000000004	: 61	, type 2 error	: 52	, type 1 error	: 9
Errors at threshold	0.6900000000000004	: 66	, type 2 error	: 59	, type 1 error	: 7
Errors at threshold	0.7000000000000004	: 66	, type 2 error	: 60	, type 1 error	: 6
Errors at threshold	0.7100000000000004	: 71	, type 2 error	: 65	, type 1 error	: 6
Errors at threshold	0.7200000000000004	: 72	, type 2 error	: 66	, type 1 error	: 6
Errors at threshold	0.7300000000000004	: 74	, type 2 error	: 70	, type 1 error	: 4
Errors at threshold	0.7400000000000004	: 73	, type 2 error	: 70	, type 1 error	: 3
Errors at threshold	0.7500000000000004	: 76	, type 2 error	: 74	, type 1 error	: 2
Errors at threshold	0.7600000000000005	: 80	, type 2 error	: 78	, type 1 error	: 2
Errors at threshold	0.7700000000000005	: 84	, type 2 error	: 82	, type 1 error	: 2
Errors at threshold	0.7800000000000005	: 86	, type 2 error	: 84	, type 1 error	: 2
Errors at threshold	0.7900000000000005	: 84	, type 2 error	: 84	, type 1 error	: 0
Errors at threshold	0.8000000000000005	: 84	, type 2 error	: 84	, type 1 error	: 0

In [53]: *# Now we decide the threshold to be 0.35 as it has lower Type 2 error and considerably acceptable Total error*
While detecting a disease we Type 2 error says that he/she doesn't have a disease in fact when he/she has the disease
Hence we look for a small Type 2 error

```
y_pred1 = np.where(y_pred_prob[:, 1] > 0.35, 1, 0)

cfm = confusion_matrix(y_test, y_pred1)
print(cfm)
print(classification_report(y_test, y_pred1))
print("Accuracy is", round(accuracy_score(y_test, y_pred1)*100, 2), "%")
```

```
[[30 41]
 [ 6 79]]
```

	precision	recall	f1-score	support
0	0.83	0.42	0.56	71
1	0.66	0.93	0.77	85
accuracy			0.70	156
macro avg	0.75	0.68	0.67	156
weighted avg	0.74	0.70	0.68	156

Accuracy is 69.87 %

A slight increase in the accuracy but not worth considering

Step 11:

Performing SVC

```
In [52]: model = SVC(C=1, kernel='rbf') # i.e cost = 1

model.fit(x_train, y_train)

y_pred = model.predict(x_test)

print(cfm)
print(classification_report(y_test, y_pred))
print("Accuracy is", round(accuracy_score(y_test, y_pred)*100, 2) , "%")
```

```
[[30 41]
 [ 6 79]]
```

	precision	recall	f1-score	support
0	0.78	0.51	0.62	71
1	0.68	0.88	0.77	85
accuracy			0.71	156
macro avg	0.73	0.69	0.69	156
weighted avg	0.73	0.71	0.70	156

Accuracy is 71.15 %

Step 12:

Applying Random Forest

```
In [55]: rf = RandomForestClassifier(n_estimators=500, criterion='entropy', min_samples_split=4)
rf.fit(x_train, y_train)
ypred_rf = rf.predict(x_test)
confusion_matrix(y_test, ypred_rf)
print("Accuracy is", round(accuracy_score(y_test, ypred_rf)*100, 2) , "%")
```

Accuracy is 83.97 %

Conclusion:

Out of all the models we applied, Random Forest gave us the best accuracy. Hence we find this model suitable for our data and will use this for further evaluation if new data comes in.

In []: